

Word Embedding

Own Sentence

```
sentences = [  
    'i love my dog',  
    'I, love my cat',  
    'You love my dog!'  
]
```

Json

```
import json  
  
with open("/tmp/sarcasm.json", 'r') as f:  
    datastore = json.load(f)  
  
sentences = []  
labels = []  
urls = []  
for item in datastore:  
    sentences.append(item['headline'])  
    labels.append(item['is_sarcastic'])  
    urls.append(item['article_link'])
```

CSV

```
sentences = []  
labels = []  
with open("/tmp/bbc-text.csv", 'r') as csvfile:  
    reader = csv.reader(csvfile, delimiter=',')  
    next(reader)  
    for row in reader:
```

cleaning for stopwords

```
for row in reader:  
    labels.append(row[0])  
    sentence = row[1]  
    for word in stopwords:  
        token = " " + word + " "  
        sentence = sentence.replace(token, " ")  
        sentence = sentence.replace(" ", " ")  
    sentences.append(sentence)
```

“sentences”: list of strings
“labels”: list of numbers

split “sentences” and “labels” to train-test sets

training_sentences testing_sentences
training_labels testing_labels

tokenizer and pad_sequences (only work with sentences and keep the labels as is)

```
tokenizer = Tokenizer(num_words=vocab_size, oov_token=oov_tok)  
tokenizer.fit_on_texts(training_sentences)
```

```
word_index = tokenizer.word_index
```

```
training_sequences = tokenizer.texts_to_sequences(training_sentences)  
training_padded = pad_sequences(training_sequences, maxlen=max_length,  
                                padding=padding_type, truncating=trunc_type)
```

```
testing_sequences = tokenizer.texts_to_sequences(testing_sentences)  
testing_padded = pad_sequences(testing_sequences, maxlen=max_length,  
                               padding=padding_type, truncating=trunc_type)
```

Hyperparameters
** tokenizer
*num_words = vocab_size (given)

**pad_sequences
*sequence
*maxlen
*padding
*truncating

If the labels are given as list of strings:

```
label_tokenizer = Tokenizer()  
label_tokenizer.fit_on_texts(labels)
```

```
training_label_seq = label_tokenizer.texts_to_sequences(train_labels)  
validation_label_seq = label_tokenizer.texts_to_sequences(validation_labels)
```

Convert to numpy arrays the ff:

“training_sequences” “testing_sequences”
“training_labels” “testing_labels”

“train_padded”: array “test_padded”: array
“train_labels”: array of list of numbers - shape (no. of samples, 1) “test_labels”: array of list of numbers - shape (no. of samples, 1)

Build Model

```
model = tf.keras.Sequential([
    tf.keras.layers.Embedding(vocab_size, embedding_dim, input_length=max_length),

    tf.keras.layers.GlobalAveragePooling1D(),
    # alternatives
    # tf.keras.layers.GlobalMaxPooling1D()
    # tf.keras.layers.Flatten()

    tf.keras.layers.Dense(24, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

```
model.compile(loss='binary_crossentropy',
              optimizer='adam',
              metrics=['accuracy'])
```

```
model.summary()
```

```
num_epochs = 30
history = model.fit(training_padded,
                    training_labels,
                    epochs=num_epochs,
                    validation_data=(testing_padded, testing_labels),
                    verbose=2)
```

Plot Loss and Accuracy

```
import matplotlib.pyplot as plt

def plot_graphs(history, string):
    plt.plot(history.history[string])
    plt.plot(history.history['val_'+string])
    plt.xlabel("Epochs")
    plt.ylabel(string)
    plt.legend([string, 'val_'+string])
    plt.show()

plot_graphs(history, "accuracy")
plot_graphs(history, "loss")
```

Creating tsv data

```
reverse_word_index = dict([(value, key) for (key, value) in word_index.items()])  
  
def decode_sentence(text):  
    return ' '.join([reverse_word_index.get(i, '?') for i in text])
```

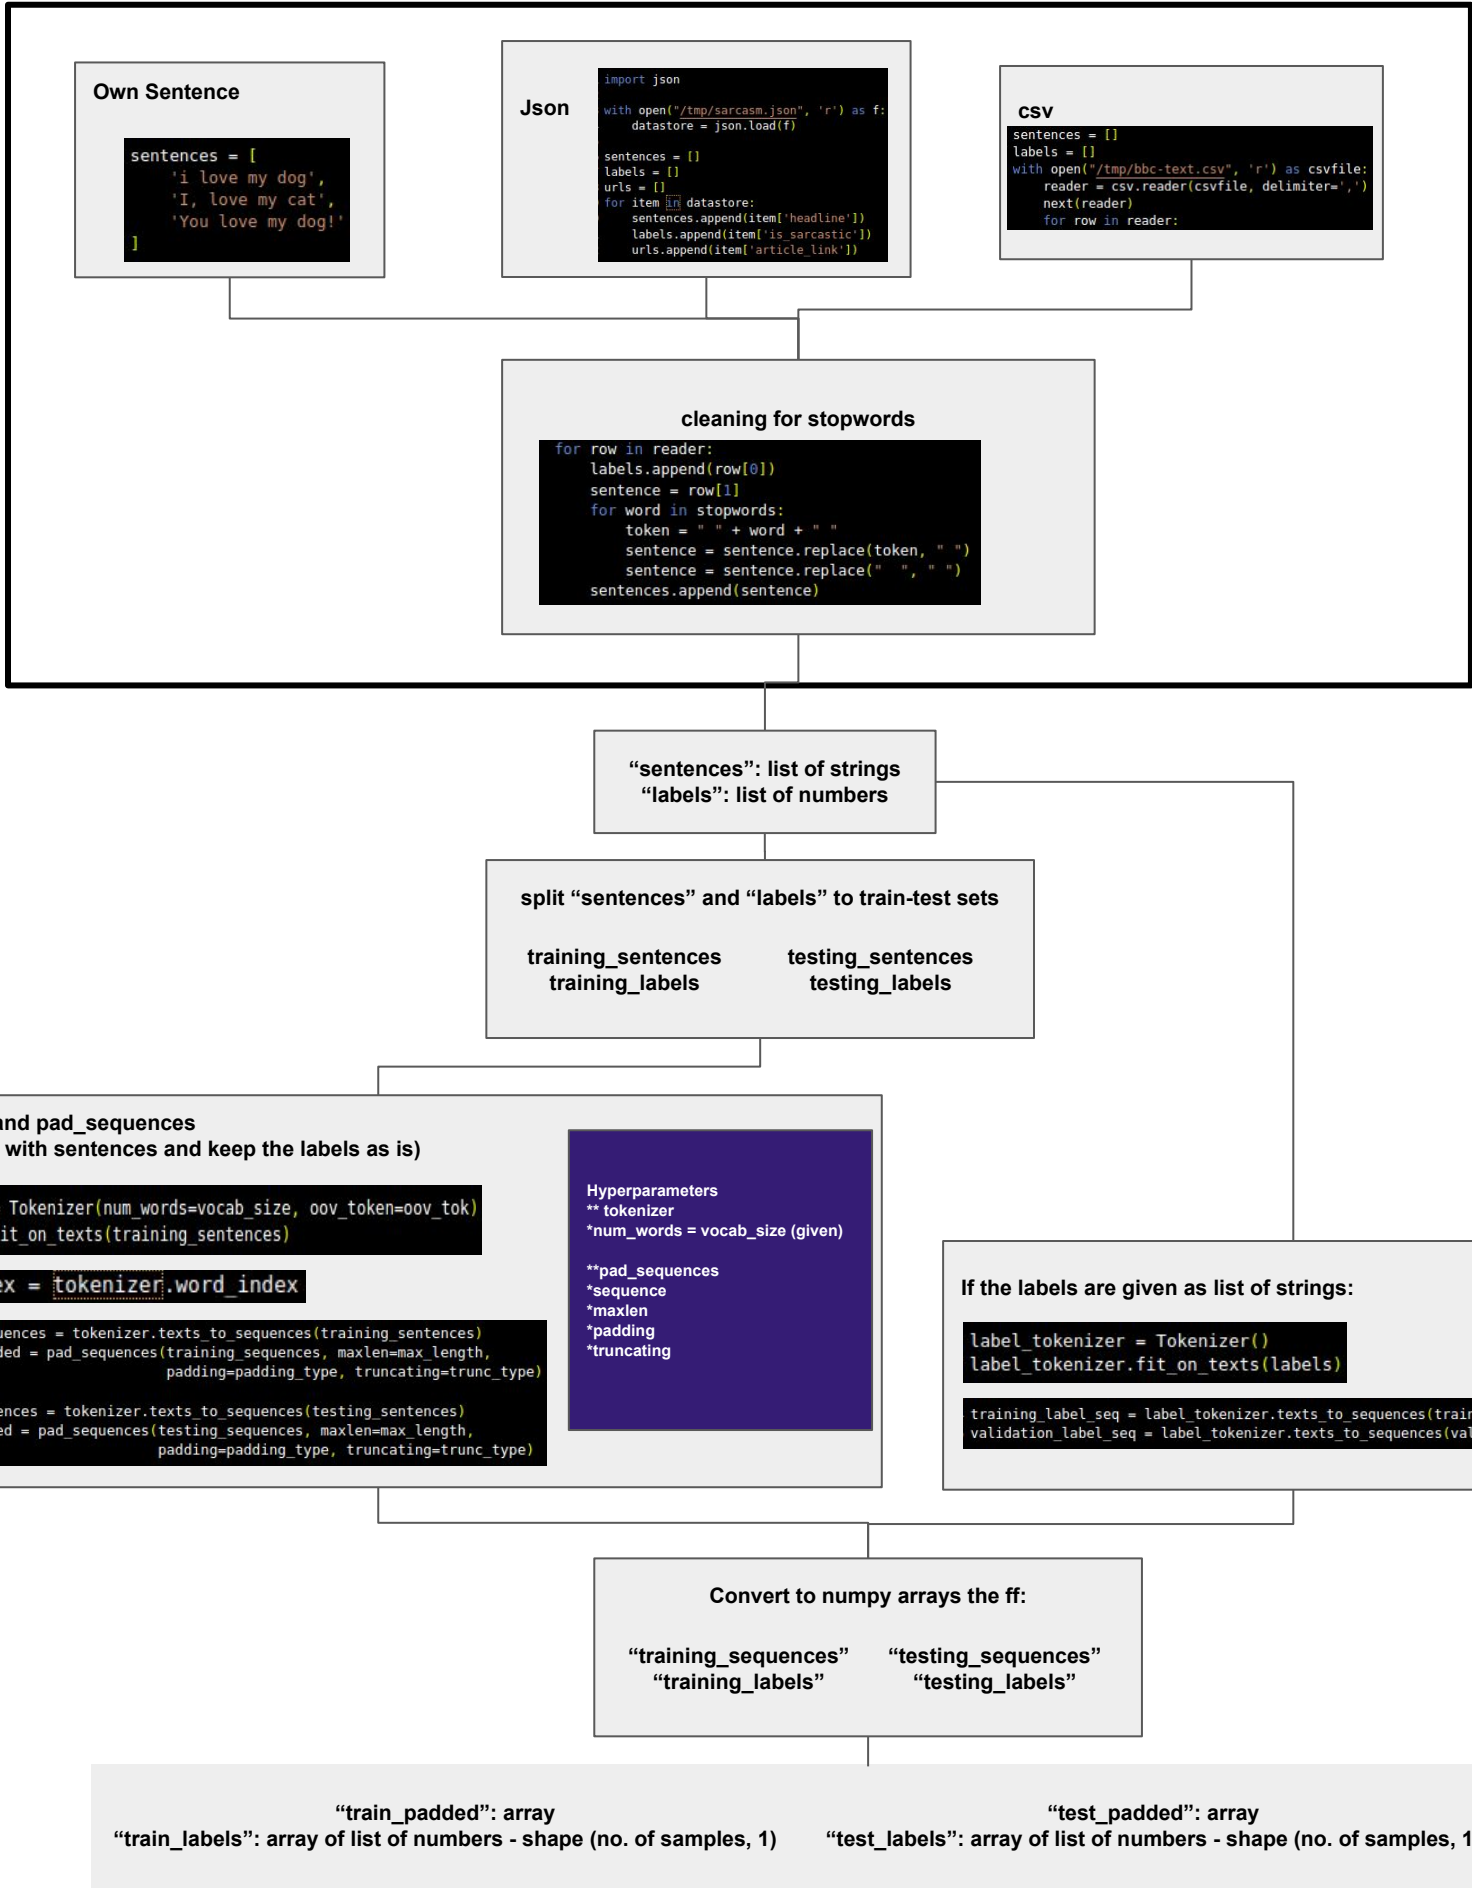
```
e = model.layers[0]  
weights = e.get_weights()[0]  
print(weights.shape) # shape: (vocab_size, embedding_dim)
```

```
import io  
  
out_v = io.open('vecs.tsv', 'w', encoding='utf-8')  
out_m = io.open('meta.tsv', 'w', encoding='utf-8')  
  
for word_num in range(1, vocab_size):  
  
    word = reverse_word_index[word_num]  
    embeddings = weights[word_num]  
  
    out_m.write(word + "\n")  
    out_v.write('\t'.join([str(x) for x in embeddings]) + "\n")  
  
out_v.close()  
out_m.close()
```

*Download files from colab

```
try:  
    from google.colab import files  
except ImportError:  
    pass  
else:  
    files.download('vecs.tsv')  
    files.download('meta.tsv')
```

Classification without Transfer Learning



Embedding Layer only and GlobalAveragePooling1D

```
model = tf.keras.Sequential([
    tf.keras.layers.Embedding(vocab_size, embedding_dim, input_length=max_length),

    tf.keras.layers.GlobalAveragePooling1D(),
    # alternatives
    # tf.keras.layers.GlobalMaxPooling1D()
    # tf.keras.layers.Flatten()

    tf.keras.layers.Dense(24, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

Bidirectional and Stacked LSTMs

```
model = tf.keras.Sequential([
    tf.keras.layers.Embedding(tokenizer.vocab_size, 64),
    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(64, return_sequences=True)),
    tf.keras.layers.Bidirectional(tf.keras.layers.LSTM(32)),
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

Conv1D

```
model = tf.keras.Sequential([
    tf.keras.layers.Embedding(vocab_size, embedding_dim, input_length=max_length),
    tf.keras.layers.Conv1D(128, 5, activation='relu'),
    tf.keras.layers.GlobalMaxPooling1D(),
    tf.keras.layers.Dense(24, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

GRU

```
model = tf.keras.Sequential([
    tf.keras.layers.Embedding(vocab_size, embedding_dim, input_length=max_length),
    tf.keras.layers.Bidirectional(tf.keras.layers.GRU(32)),
    tf.keras.layers.Dense(6, activation='relu'),
    tf.keras.layers.Dense(1, activation='sigmoid')
])
```

Stacking LSTM and Conv1D

```
model = tf.keras.Sequential([
    tf.keras.layers.Embedding(vocab_size, embedding_dim, input_length=max_length),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Conv1D(64, 5, activation="relu"),
    tf.keras.layers.MaxPooling1D(pool_size=4),
    tf.keras.layers.LSTM(64),
    tf.keras.layers.Dense(1, activation="sigmoid")
])
```

Model Compile, Summary, Fit

```
model.compile(loss='binary_crossentropy',
              optimizer='adam',
              metrics=['accuracy'])
```

```
model.summary()
```

```
num_epochs = 30
history = model.fit(training_padded,
                    training_labels,
                    epochs=num_epochs,
                    validation_data=(testing_padded, testing_labels),
                    verbose=2)
```

Model Predict

```
sentence = ["granny starting to fear spiders in the garden might be real",
            "game of thrones season finale showing this sunday night"]

sequences = tokenizer.texts_to_sequences(sentence)

padded = pad_sequences(sequences, maxlen=max_length,
                      padding=padding_type, truncating=trunc_type)

print(model.predict_classes(padded))
```

Plot Loss and Accuracy

```
import matplotlib.pyplot as plt

def plot_graphs(history, string):
    plt.plot(history.history[string])
    plt.plot(history.history['val_'+string])
    plt.xlabel("Epochs")
    plt.ylabel(string)
    plt.legend([string, 'val_'+string])
    plt.show()

plot_graphs(history, "accuracy")
plot_graphs(history, "loss")
```

Classification with Transfer Learning - week 3

Own Sentence

```
sentences = [  
    'i love my dog',  
    'I, love my cat',  
    'You love my dog!'  
]
```

Json

```
import json  
  
with open("/tmp/sarcasm.json", 'r') as f:  
    datastore = json.load(f)  
  
sentences = []  
labels = []  
urls = []  
for item in datastore:  
    sentences.append(item['headline'])  
    labels.append(item['is_sarcastic'])  
    urls.append(item['article_link'])
```

CSV

```
sentences = []  
labels = []  
with open("/tmp/bbc-text.csv", 'r') as csvfile:  
    reader = csv.reader(csvfile, delimiter=',')  
    next(reader)  
    for row in reader:
```

cleaning for stopwords

```
for row in reader:  
    labels.append(row[0])  
    sentence = row[1]  
    for word in stopwords:  
        token = " " + word + " "  
        sentence = sentence.replace(token, " ")  
        sentence = sentence.replace(" ", " ")  
    sentences.append(sentence)
```

***This may vary a lot**
Also see

```
import tensorflow_datasets as tfds  
imdb, info = tfds.load("imdb_reviews", with_info=True, as_supervised=True)
```

“sentences”: list of strings
“labels”: list of numbers

split “sentences” and “labels” to train-test sets

training_sentences testing_sentences
training_labels testing_labels

tokenizer and pad_sequences (only work with sentences and keep the labels as is)

```
tokenizer = Tokenizer(num_words=vocab_size, oov_token=oov_tok)  
tokenizer.fit_on_texts(training_sentences)
```

```
word_index = tokenizer.word_index
```

```
training_sequences = tokenizer.texts_to_sequences(training_sentences)  
training_padded = pad_sequences(training_sequences, maxlen=max_length,  
                                padding=padding_type, truncating=trunc_type)
```

```
testing_sequences = tokenizer.texts_to_sequences(testing_sentences)  
testing_padded = pad_sequences(testing_sequences, maxlen=max_length,  
                               padding=padding_type, truncating=trunc_type)
```

Hyperparameters
** tokenizer
* num_words = vocab_size (given)

** pad_sequences
* sequence
* maxlen
* padding
* truncating

If the labels are given as list of strings:

```
label_tokenizer = Tokenizer()  
label_tokenizer.fit_on_texts(labels)
```

```
training_label_seq = label_tokenizer.texts_to_sequences(train_labels)  
validation_label_seq = label_tokenizer.texts_to_sequences(validation_labels)
```

Convert to numpy arrays the ff:

“training_sequences” “testing_sequences”
“training_labels” “testing_labels”

“train_padded”: array “test_padded”: array
“train_labels”: array of list of numbers - shape (no. of samples, 1) “test_labels”: array of list of numbers - shape (no. of samples, 1)

Get Embeddings from GloVe

```
# Note this is the 100 dimension version of GloVe from Stanford
# I unzipped and hosted it on my site to make this notebook easier
!wget --no-check-certificate \
  https://storage.googleapis.com/laurencemoroney-blog.appspot.com/glove.6B.100d.txt \
  -O /tmp/glove.6B.100d.txt

embeddings_index = {};
with open('/tmp/glove.6B.100d.txt') as f:
    for line in f:
        values = line.split();
        word = values[0];
        coefs = np.asarray(values[1:], dtype='float32');
        embeddings_index[word] = coefs;

embeddings_matrix = np.zeros((vocab_size+1, embedding_dim));
for word, i in word_index.items():
    embedding_vector = embeddings_index.get(word);
    if embedding_vector is not None:
        embeddings_matrix[i] = embedding_vector;
```

Model (See changes in Embedding layer)

```
model = tf.keras.Sequential([
    tf.keras.layers.Embedding(vocab_size+1, embedding_dim, input_length=max_length,
                             weights=[embeddings_matrix], trainable=False),
    tf.keras.layers.Dropout(0.2),
    tf.keras.layers.Conv1D(64, 5, activation="relu"),
    tf.keras.layers.MaxPooling1D(pool_size=4),
    tf.keras.layers.LSTM(64),
    tf.keras.layers.Dense(1, activation="sigmoid")
])
```

Model Compile, Summary, Fit

```
model.compile(loss='binary_crossentropy',
              optimizer='adam',
              metrics=['accuracy'])
```

```
model.summary()
```

```
num_epochs = 30
history = model.fit(training_padded,
                    training_labels,
                    epochs=num_epochs,
                    validation_data=(testing_padded, testing_labels),
                    verbose=2)
```

Plot Loss and Accuracy

```
import matplotlib.pyplot as plt

def plot_graphs(history, string):
    plt.plot(history.history[string])
    plt.plot(history.history['val_'+string])
    plt.xlabel("Epochs")
    plt.ylabel(string)
    plt.legend([string, 'val_'+string])
    plt.show()

plot_graphs(history, "accuracy")
plot_graphs(history, "loss")
```

Model Predict

```
sentence = ["granny starting to fear spiders in the garden might be real",
            "game of thrones season finale showing this sunday night"]

sequences = tokenizer.texts_to_sequences(sentence)

padded = pad_sequences(sequences, maxlen=max_length,
                      padding=padding_type, truncating=trunc_type)

print(model.predict_classes(padded))
```

Text(Song/Poetry) Generation

```
from tensorflow.keras.preprocessing.sequence import pad_sequences
from tensorflow.keras.layers import Embedding, LSTM, Dense, Dropout, Bidirectional
from tensorflow.keras.preprocessing.text import Tokenizer
from tensorflow.keras.models import Sequential
from tensorflow.keras.optimizers import Adam
### YOUR CODE HERE
# Figure out how to import regularizers
from tensorflow.keras.regularizers import L1, L2, L1L2
###
import tensorflow.keras.utils as ku
import numpy as np

from tensorflow import keras
```

Get the data

```
!wget --no-check-certificate \
  https://storage.googleapis.com/laurencemoroney-blog.appspot.com/sonnets.txt \
  -O /tmp/sonnets.txt

data = open('/tmp/sonnets.txt').read()
corpus = data.lower().split("\n")
```

Tokenize the corpus

```
tokenizer = Tokenizer()
tokenizer.fit_on_texts(corpus)
total_words = len(tokenizer.word_index) + 1
```

N-gram and Input Sequences

```
# create input sequences using list of tokens
input_sequences = []
for line in corpus:
    token_list = tokenizer.texts_to_sequences([line])[0]
    for i in range(1, len(token_list)):
        n_gram_sequence = token_list[:i+1]
        input_sequences.append(n_gram_sequence)
```

Pad Sequences

```
max_sequence_len = max([len(x) for x in input_sequences])
input_sequences = np.array(pad_sequences(input_sequences, maxlen=max_sequence_len, padding='pre'))
```

Inputs and Labels

```
xs, label = input_sequences[:, :-1], input_sequences[:, -1]

label = ku.to_categorical(label, num_classes=total_words)
```


Model

```
model = keras.models.Sequential([
    keras.layers.Embedding(total_words, 64, input_length = max_sequence_len - 1),
    keras.layers.Bidirectional(keras.layers.LSTM(20, return_sequences = True)),
    keras.layers.Dropout(0.2),
    keras.layers.Bidirectional(keras.layers.LSTM(20)),
    keras.layers.Dense(256, kernel_regularizer=L2(0.01), activation="relu"),
    keras.layers.Dense(total_words, activation="softmax")
])

model.compile(loss = "categorical_crossentropy", optimizer = "adam", metrics = ["accuracy"])
print(model.summary())
```

```
history = model.fit(xs, label, epochs=100, verbose=1)
```

Plot Accuracy and Loss

```
import matplotlib.pyplot as plt
acc = history.history['accuracy']
loss = history.history['loss']

epochs = range(len(acc))

plt.plot(epochs, acc, 'b', label='Training accuracy')
plt.title('Training accuracy')

plt.figure()

plt.plot(epochs, loss, 'b', label='Training Loss')
plt.title('Training loss')
plt.legend()

plt.show()
```

Predict - Generate Sonnet/Songs/Poetry

```
seed_text = "Help me Obi Wan Kenobi, you're my only hope"
next_words = 100

for _ in range(next_words):
    token_list = tokenizer.texts_to_sequences([seed_text])[0]
    token_list = pad_sequences([token_list], maxlen=max_sequence_len-1, padding='pre')
    predicted = model.predict_classes(token_list, verbose=0)
    output_word = ""
    for word, index in tokenizer.word_index.items():
        if index == predicted:
            output_word = word
            break
    seed_text += " " + output_word
print(seed_text)
```